

Machine Learning on Graphs: Foundations and Applications

EXERCISE 2

Group 29

Rishang Yadav Risa Niranjan Chaudhari Bhoomi Dhawan
485041 485510 485556

-
1. **Problem 1:** For $k > 1$, prove that the folklore k -WL is at least as expressive as the δ - k -WL in distinguishing non-isomorphic graphs.

Ans: We prove that whenever two k -tuples s, t receive the same stable colour under the folklore k -WL, then they also receive the same stable colour under the δ - k -WL. The proof proceeds by induction using unrolling trees.

Unrolling trees. For a WL-type algorithm \mathcal{A} and a tuple $v \in V(G)^k$, let $\text{UNR}^{\mathcal{A}}[G, v, i]$ be the depth- i unrolling tree. Standard WL theory gives:

$$C_{i,\mathcal{A}}^k(v) = C_{i,\mathcal{A}}^k(w) \iff \text{UNR}^{\mathcal{A}}[G, v, i] \cong \text{UNR}^{\mathcal{A}}[G, w, i].$$

Thus two tuples have the same stable WL-colour iff all their unrolling trees of all depths are isomorphic.

Suppose

$$\text{UNR}^{\text{folklore}}[G, s, q] \cong \text{UNR}^{\text{folklore}}[G, t, q]$$

for some q exceeding the folklore WL stabilization time. We prove by induction on i that

$$\text{UNR}^{\delta}[G, s, i] \cong \text{UNR}^{\delta}[G, t, i].$$

Base case $i = 0$. Zero depth unrollings consist only of the atomic type of the tuple. Since the folklore unrolling roots are isomorphic, s and t have the same atomic type, hence

$$\text{UNR}^{\delta}[G, s, 0] \cong \text{UNR}^{\delta}[G, t, 0].$$

Inductive step. Assuming the claim holds for depth $i - 1$, we now show it for depth i .

For each position $j \in \{1, \dots, k\}$ define

$$X_j = \{\varphi_j(s, w) : w \in V(G)\}, \quad Y_j = \{\varphi_j(t, w) : w \in V(G)\},$$

i.e. all tuples obtained by replacing the j -th component. In the δ -unrolling, the children of s at position j are exactly the depth- $(i - 1)$ unrollings of elements of X_j , with edge label $(j, \text{adj}(s_j, w))$, and similarly for Y_j .

Folklore vectors determine adjacency. In folklore k -WL, the neighbour of s obtained by replacing v_j with w contributes the vector

$$(C(\varphi_1(s, w)), \dots, C(\varphi_k(s, w))).$$

The colour $C(\varphi_j(s, w))$ contains the atomic type of $\varphi_j(s, w)$, which encodes whether w is adjacent to s_j . Hence the adjacency bit required by δ -WL is already determined by the folklore colour vector.

Matching children. Since the folklore unrolling trees of s and t are isomorphic at depth q , the multisets of folklore colour-vectors obtained from s and t coincide. Consequently, for each position j and each folklore colour-vector class \mathbf{c} , the number of elements of X_j of class \mathbf{c} equals that of Y_j . Because \mathbf{c} determines the adjacency bit, the numbers of local and global δ -children in each class match.

Thus for each j , we can have a bijection

$$\sigma_j : X_j \rightarrow Y_j$$

preserving adjacency bits and folklore colours.

Applying the induction hypothesis. If $x \in X_j$ and $\sigma_j(x) \in Y_j$ have the same folklore colour at depth $i - 1$, then their folklore unrollings of depth $i - 1$ are isomorphic. By the induction hypothesis,

$$\text{UNR}^\delta[G, x, i - 1] \cong \text{UNR}^\delta[G, \sigma_j(x), i - 1].$$

Conclusion of the step. The children of s and t at position j are exactly the trees rooted at X_j and Y_j with their adjacency labels. Since σ_j preserves adjacency and pairs each child with an isomorphic $(i - 1)$ -subtree, we obtain an isomorphism

$$\text{UNR}^\delta[G, s, i] \cong \text{UNR}^\delta[G, t, i].$$

Final conclusion. By induction, all depths match. Hence the stable colours satisfy

$$C_{\infty, \delta}^k(s) = C_{\infty, \delta}^k(t).$$

Therefore **folklore k -WL is at least as expressive as δ - k -WL for all $k > 1$** .



2. **Problem 2:** For $k > 1$, prove that the δ - k -WL is at least as expressive as the local- k -WL in distinguishing non-isomorphic graphs. (2 Pt.)

Ans: The proof relies on the equivalence between stable coloring and the isomorphism of the corresponding unrolling trees, $\text{UNR}^{\mathcal{A}}[G, v, i]$.

Induction on Depth i

Let q be a sufficiently large depth. The assumption $C_{\infty, L,+}^k(s) = C_{\infty, L,+}^k(t)$ implies $\text{UNR}^{L,+}[G, s, q] \cong \text{UNR}^{L,+}[G, t, q]$. We induct on i to show $\text{UNR}^{\delta}[G, s, i] \cong \text{UNR}^{\delta}[G, t, i]$ for all i .

Base Case $i = 0$: The isomorphism of the roots of $\text{UNR}^{L,+}[G, s, q]$ and $\text{UNR}^{L,+}[G, t, q]$ implies they have the same atomic type. Since $\text{UNR}^{\delta}[G, v, 0]$ is a single node with the atomic type, $\text{UNR}^{\delta}[G, s, 0] \cong \text{UNR}^{\delta}[G, t, 0]$.

Inductive Step i : Assume $\text{UNR}^{\delta}[G, x, i - 1] \cong \text{UNR}^{\delta}[G, \sigma_j(x), i - 1]$ for all j -neighbors x and their image $\sigma_j(x)$.

1. **Matching via $L - \text{WL}^+$:** The isomorphism $\text{UNR}^{L,+}[G, s, q] \cong \text{UNR}^{L,+}[G, t, q]$ ensures that for any stable color class C and position j :

$$|X_j \cap C| = |Y_j \cap C|,$$

and, crucially, the counts of **local** (adjacent) and **global** (non-adjacent) neighbors in C are also individually equal for s and t .

2. **Bijection σ_j :** This count equality allows us to construct a bijection $\sigma_j : X_j \rightarrow Y_j$ that pairs up j -neighbors x and $\sigma_j(x)$ such that they are in the same stable L - WL^+ color class and share the same adjacency type (local/global).
3. **Isomorphism δ -UNR:** Since σ_j preserves the L - WL^+ color, by the induction hypothesis, the δ -unrolling subtrees are isomorphic: $\text{UNR}^{\delta}[G, x, i-1] \cong \text{UNR}^{\delta}[G, \sigma_j(x), i-1]$. The δ -WL uses the j -label and the adjacency type $adj(v, w) \in \{0, 1\}$ (corresponding to local/global) as its edge label. Since σ_j preserves both the subtree isomorphism and the adjacency label, we obtain the overall isomorphism for depth i :

$$\text{UNR}^{\delta}[G, s, i] \cong \text{UNR}^{\delta}[G, t, i].$$

Conclusion By induction, $C_{\infty, \delta}^k(s) = C_{\infty, \delta}^k(t)$.



3. **Problem 3:** Show that the local k -WL⁺ can be computed in the same asymptotic running time as the local k -WL.

Ans: We need to show that both algorithms have the same asymptotic complexity of $O(n^k \cdot k \cdot d)$ per iteration, where $n = |V(G)|$ is the number of vertices, k is the tuple dimension, and d is the maximum degree.

Aggregation functions. For local k -WL:

$$M_{t,L}(v) = \left(\{\{C_t^{k,L}(\varphi_1(v, w)) \mid w \in N(v_1)\}\}, \dots, \{\{C_t^{k,L}(\varphi_k(v, w)) \mid w \in N(v_k)\}\} \right)$$

For local k -WL⁺:

$$M_{t,L,+}(v) = \left(\{\{(C_t^{k,L,+}(\varphi_j(v, w)), \#_t^j(v, \varphi_j(v, w))) \mid w \in N(v_j)\}\}_{j=1}^k \right)$$

where

$$\#_t^j(v, x) = |\{w : w \sim_j v, C_t^{k,L,+}(w) = C_t^{k,L,+}(x)\}|.$$

Running time analysis for local k -WL. For each k -tuple $v \in V(G)^k$ at iteration t :

- For each position $j \in [k]$, iterate over the local neighbors $N(v_j)$.
- For each $w \in N(v_j)$, look up the color $C_t^{k,L}(\varphi_j(v, w))$.
- Collect these colors in a multiset.

Cost per tuple: $O(k \cdot d)$.

Total cost per iteration: $O(n^k \cdot k \cdot d)$.

Running time analysis for local k -WL⁺. The $\#$ function can be computed efficiently using hash maps without changing the asymptotic complexity.

For each k -tuple $v \in V(G)^k$ at iteration t :

First pass – count colors by position:

- For each position $j \in [k]$, create a hash map $\text{count}_j : \mathbb{N} \rightarrow \mathbb{N}$.
- For each $w \in N(v_j)$, compute $x = \varphi_j(v, w)$ and increment $\text{count}_j[C_t^{k,L,+}(x)]$.
- Time per position: $O(|N(v_j)|) = O(d)$.

Second pass – build multiset with counts:

- For each position $j \in [k]$ and each $w \in N(v_j)$:
 - Compute $x = \varphi_j(v, w)$.
 - Look up $\#_t^j(v, x) = \text{count}_j[C_t^{k,L,+}(x)]$ (constant time hash lookup).
 - Add pair $(C_t^{k,L,+}(x), \#_t^j(v, x))$ to the multiset.
- Time per position: $O(|N(v_j)|) = O(d)$.

Cost per tuple: $O(k \cdot d) + O(k \cdot d) = O(k \cdot d)$.

Total cost per iteration: $O(n^k \cdot k \cdot d)$.

The $\#$ function counts *all* j -neighbors (local and global) with a given color, but we only iterate over *local* neighbors in $M_{t,L,+}$.

For a fixed tuple v and position j , the set of all j -neighbors is fixed and can be processed once per iteration.

Counting colors via hash maps requires only $O(d)$ time per position, the same as collecting the colors themselves.

The additional bookkeeping (counting and lookup) adds only constant factor overhead per neighbor.

Conclusion. Both local k -WL and local k -WL⁺ have the same asymptotic running time of

$$O(n^k \cdot k \cdot d) \text{ per iteration.}$$



4. **Problem 4:** Prove that the local k -WL induces a hierarchy. That is, for $k \geq 1$, show that there exists a pair of graphs that the local k -WL cannot distinguish but the local $(k+1)$ -WL can.

Ans: We construct, for each $k \geq 1$, a pair of non-isomorphic graphs (G_k, H_k) such that the local k -WL cannot distinguish them, but the local $(k+1)$ -WL can.

Construction of G_k and H_k . Let K denote the complete graph on $k+1$ vertices (without self-loops). The vertices of K are numbered from 0 to k . Let $E(v)$ denote the set of edges incident to v in K . Clearly, $|E(v)| = k$ for all $v \in V(K)$.

Graph G_k :

1. For the vertex set $V(G_k)$, we add:

- (v, S) for each $v \in V(K)$ and for each *even* subset S of $E(v)$ (including \emptyset).
 - Two vertices e_0, e_1 for each edge $e \in E(K)$.
2. For the edge set $E(G_k)$, we add:
- An edge $\{e_0, e_1\}$ for each $e \in E(K)$.
 - An edge between (v, S) and e_1 if $v \in e$ and $e \in S$.
 - An edge between (v, S) and e_0 if $v \in e$ and $e \notin S$.

Graph H_k : Define H_k similarly to G_k , with the following exception:

- For vertex $0 \in V(K)$, we choose all *odd* subsets of $E(0)$ instead of even subsets.

Note that $|V(G_k)| = |V(H_k)| = (k+1) \cdot 2^{k-1} + \binom{k+1}{2} \cdot 2$.

The graphs are non-isomorphic. We use the concept of a *distance-two-clique*: a set S of vertices is a distance-two-clique if the distance between any two vertices in S is exactly two.

1. There exists a distance-two-clique of size $(k+1)$ in G_k .
2. There does not exist a distance-two-clique of size $(k+1)$ in H_k .

Hence G_k and H_k are non-isomorphic.

Part 1: In G_k , consider $S = \{(0, \emptyset), (1, \emptyset), \dots, (k, \emptyset)\}$ of size $k+1$. For any $i, j \in V(K)$ with $i \neq j$, vertex (i, \emptyset) is adjacent to $\{i, j\}_0$, which is adjacent to (j, \emptyset) . Thus any two vertices in S are at distance two.

Part 2: Suppose there exists a distance-two-clique $\{(0, S_0), (1, S_1), \dots, (k, S_k)\}$ in H_k , where $S_i \subseteq E(i)$.

Computing the parity-sum of $|S_0|, \dots, |S_k|$: Since S_0 is odd and all others are even, the total parity is 1.

Computing edge-by-edge: For each edge $\{i, j\} \in E(K)$, since (i, S_i) and (j, S_j) are at distance two, either both S_i and S_j contain $\{i, j\}$ or neither does. Thus the parity contribution of $\{i, j\}$ to the sum is 0.

Since each edge contributes 0, the total parity-sum must be 0, a contradiction. Hence no such distance-two-clique exists in H_k .

The local k -WL cannot distinguish G_k and H_k . This follows from classical WL theory. The graphs G_k and H_k are the standard Cai-Fürer-Immerman (CFI)

construction, which is known to fool the k -WL. Since the local k -WL only uses local neighbors, which is a subset of the information used by the standard k -WL, it follows that:

$$\text{local } k\text{-WL} \preceq k\text{-WL}.$$

Therefore, the local k -WL also cannot distinguish G_k and H_k .

The local $(k+1)$ -WL can distinguish G_k and H_k . We show that the local $(k+1)$ -WL can detect the distance-two-clique in G_k .

Consider the $(k+1)$ -tuple

$$P = ((0, \emptyset), (1, \emptyset), \dots, (k, \emptyset)) \in V(G_k)^{k+1}.$$

Claim: There is no tuple $Q \in V(H_k)^{k+1}$ such that the local $(k+1)$ -WL unrolling of P is isomorphic to the unrolling of Q .

[Proof of Claim] Suppose, for contradiction, that such a Q exists. By comparing atomic types, Q must be of the form $((0, S_0), (1, S_1), \dots, (k, S_k))$.

Consider the depth-two unrolling of P : From the root P , we can go down via two local edges labeled 1 to reach the tuple

$$P_2 = ((1, \emptyset), (1, \emptyset), (2, \emptyset), \dots, (k, \emptyset)).$$

This is possible because $(0, \emptyset)$ and $(1, \emptyset)$ are at distance two via the edge-gadget.

For the unrolling of Q to be isomorphic, vertices $(0, S_0)$ and $(1, S_1)$ must be at distance two in H_k . Repeating this argument for all pairs of positions, we find that (i, S_i) and (j, S_j) must be at distance two for all $i \neq j$.

Now consider the depth-four unrolling of P : From P , we can reach $R = ((0, \emptyset), (1, \emptyset), \dots, (k, \emptyset))$ with position 1 replaced by $(0, \emptyset)$. This shows that $(0, \emptyset)$ is at distance two from all (i, \emptyset) for $i \in \{1, \dots, k\}$.

Similarly, in H_k , vertex $(0, S_0)$ must be at distance two from all (i, S_i) for $i \in \{1, \dots, k\}$. Combined with the pairwise distances, this means $\{(0, S_0), (1, S_1), \dots, (k, S_k)\}$ forms a distance-two-clique of size $k+1$ in H_k , contradicting the lemma.

Since no such Q exists, the local $(k+1)$ -WL assigns different colors to tuples in G_k and H_k , and thus distinguishes the graphs.

Conclusion. For each $k \geq 1$, the pair (G_k, H_k) witnesses that the local $(k+1)$ -WL is strictly more powerful than the local k -WL. Therefore, the local k -WL induces a strict hierarchy:

$$\text{local 1-WL} \prec \text{local 2-WL} \prec \text{local 3-WL} \prec \dots$$



5. **Problem 5:** The random-walk kernel counts how many walks two graphs have in common. To define it, we need to introduce the direct product graph of two graphs. Let G and H be two node-labeled graphs. Then the direct product graph $G \times H$ of the two graphs is the graph with the vertex set

$$V(G \times H) := \{(g, h) \mid g \in V(G), h \in V(H), l(g) = l(h)\},$$

and the edge set

$$E(G \times H) := \{((g, h), (s, t)) \mid (g, s) \in E(G), (h, t) \in E(H)\}.$$

Now given two graphs G and H and $k > 0$, the random-walk kernel

$$k_{\text{RW}}(G, H) := \sum_{i,j=1}^{|V(G \times H)|} \left[\sum_{l=1}^k \lambda_l A(G \times H)^l \right]_{ij},$$

where $\lambda_1, \dots, \lambda_k$ with $\lambda_i > 0$ is a sequence of weights ensuring convergence, e.g., when $k = \infty$. Show that the random-walk kernel for any $k > 0$ cannot distinguish more non-isomorphic graphs than the 1-WL.

Ans: We need to prove that if the 1-WL cannot distinguish two graphs G and H , then $k_{\text{RW}}(G, G) = k_{\text{RW}}(H, H)$ for any $k > 0$. This proves that the random-walk kernel cannot distinguish more non-isomorphic graphs than the 1-WL.

Random walks and color refinement. The random-walk kernel counts walks of different lengths in the direct product graph. The crucial insight is that walks in G correspond to sequences of vertex labels and adjacencies, which are exactly the information captured by the 1-WL color refinement.

Let $C_t^1 : V(G) \rightarrow \mathbb{N}$ denote the coloring of G after t iterations of the 1-WL, and similarly for H .

$$C_0^1(v) = l(v) \quad (\text{initial vertex label})$$

$$C_{t+1}^1(v) = \text{Recolor}(C_t^1(v), \{\{C_t^1(w) \mid w \in N(v)\}\}).$$

The 1-WL assigns the same stable coloring to G and H if and only if, for all $t \geq 0$, the multisets of colors $\{\{C_t^1(v) \mid v \in V(G)\}\}$ and $\{\{C_t^1(v) \mid v \in V(H)\}\}$ are equal.

Relating walks to 1-WL colors. A walk of length l in a graph G starting at vertex v is a sequence $v = v_0, v_1, \dots, v_l$ where $(v_i, v_{i+1}) \in E(G)$ for all i .

Claim: Two vertices $v \in V(G)$ and $w \in V(H)$ have the same 1-WL color after t iterations if and only if they have the same distribution of walks of length up to t (counted by the sequence of labels encountered).

[Proof] By induction on t .

Base case $t = 0$: $C_0^1(v) = C_0^1(w)$ if and only if $l(v) = l(w)$, which means they have the same "walk of length 0" (i.e., the vertex label itself).

Inductive step: Assuming that the claim holds for $t - 1$. Then $C_t^1(v) = C_t^1(w)$ iff:

- $C_{t-1}^1(v) = C_{t-1}^1(w)$ (same walks up to length $t - 1$), and
- $\{\{C_{t-1}^1(u) \mid u \in N(v)\}\} = \{\{C_{t-1}^1(u') \mid u' \in N(w)\}\}$ (same multiset of neighbor colors).

The second condition means that for every walk of length $t - 1$ starting at a neighbor of v , there exists a corresponding walk starting at a neighbor of w with the same label sequence (and vice versa). Combined with the first condition, this means v and w have the same distribution of walks of length up to t .

Counting walks in the direct product. The entry $[A(G \times H)^l]_{(g,h),(s,t)}$ counts the number of walks of length l from (g, h) to (s, t) in the direct product graph $G \times H$.

A walk from (g, h) to (s, t) in $G \times H$ corresponds to a pair of walks:

- A walk $g = g_0, g_1, \dots, g_l = s$ in G , and
- A walk $h = h_0, h_1, \dots, h_l = t$ in H ,

such that $l(g_i) = l(h_i)$ for all $i \in \{0, \dots, l\}$ (since vertices in $G \times H$ require matching labels).

Therefore:

$$\sum_{i,j=1}^{|V(G \times H)|} [A(G \times H)^l]_{ij}$$

counts the total number of walks of length l in $G \times H$, summed over all pairs of starting and ending vertices.

1-WL indistinguishability implies equal kernel values. Assume the 1-WL cannot distinguish G and H , i.e., for all $t \geq 0$:

$$\{\{C_t^1(v) \mid v \in V(G)\}\} = \{\{C_t^1(v) \mid v \in V(H)\}\}.$$

We compute $k_{\text{RW}}(G, G)$ and $k_{\text{RW}}(H, H)$.

For $G \times G$:

- $V(G \times G) = \{(g, g') \mid g, g' \in V(G), l(g) = l(g')\}$.
- A walk of length l in $G \times G$ from (g_1, g_2) to (s_1, s_2) corresponds to two walks in G : one from g_1 to s_1 and one from g_2 to s_2 , where the label sequences match at each step.

Similarly for $H \times H$.

The number of walks of length l in $G \times G$ depends only on:

1. The number of vertices of each label (given by C_0^1).
2. The distribution of walks of each length up to l from vertices of each color (given by C_t^1 for $t \leq l$).

Since the 1-WL assigns the same color distribution to G and H at all iterations, the graphs have:

- The same number of vertices with each label.
- The same distribution of walks of each length from vertices of each 1-WL color.

Therefore, for each l :

$$\sum_{i,j=1}^{|V(G \times G)|} [A(G \times G)^l]_{ij} = \sum_{i,j=1}^{|V(H \times H)|} [A(H \times H)^l]_{ij}.$$

Conclusion. Since the above equality holds for all l , and the kernel is a weighted sum over l :

$$k_{\text{RW}}(G, G) = \sum_{l=1}^k \lambda_l \sum_{i,j} [A(G \times G)^l]_{ij} = \sum_{l=1}^k \lambda_l \sum_{i,j} [A(H \times H)^l]_{ij} = k_{\text{RW}}(H, H).$$

Therefore, if the 1-WL cannot distinguish G and H , then $k_{\text{RW}}(G, G) = k_{\text{RW}}(H, H)$, which means the random-walk kernel (as a graph feature for classification) cannot distinguish them either.

Hence, the random-walk kernel for any $k > 0$ cannot distinguish more non-isomorphic graphs than the 1-WL.



6. **Problem 6:** Given an n -order graph G of maximum degree d , devise an efficient polynomial-time algorithm (in n and constant d) for computing the feature map of the Graphlet kernel for $k = 4$. Simply enumerating all 4-Graphlets and then counting them is not allowed.

Ans: We devise an algorithm that computes the feature map for 4-node graphlets (connected subgraphs on 4 vertices) in time $O(n \cdot d^3)$, which is polynomial in n when d is constant.

Overview of 4-node connected graphlets. There are exactly 6 non-isomorphic connected graphs on 4 vertices:

1. P_4 : Path with 3 edges (1-2-3-4)
2. C_4 : 4-cycle (square)
3. $K_{1,3}$: Star with 3 edges (center connected to 3 leaves)
4. $P_3 + e$: Path with one chord (triangle + tail)
5. $K_4 - e$: Complete graph minus one edge (5 edges)
6. K_4 : Complete graph (6 edges)

Key idea: Avoid explicit enumeration. Instead of enumerating all $\binom{n}{4}$ possible 4-vertex subsets (which would be $O(n^4)$), we exploit the bounded degree to enumerate only **relevant** 4-vertex sets—those that actually form connected subgraphs.

Strategy: For each vertex v , consider all connected 4-graphlets containing v by exploring the local neighborhood up to distance 3 from v . Due to bounded degree d , this neighborhood has size at most $O(d^3)$.

Algorithm: Efficient 4-Graphlet Counting

- 1: **Input:** Graph $G = (V, E)$ with $|V| = n$ and maximum degree d
- 2: **Output:** Feature vector $\phi(G) \in \mathbb{R}^6$ counting each graphlet type
- 3: Initialize $\phi(G) = [0, 0, 0, 0, 0, 0]$ (counts for the 6 graphlet types)
- 4:
- 5: **for** each vertex $v \in V(G)$ **do**
- 6: /* Explore the local neighborhood of v */
- 7: $N_1(v) \leftarrow N(v)$ /* 1-hop neighbors */
- 8: $N_2(v) \leftarrow \{w \mid \exists u \in N_1(v), w \in N(u), w \neq v\}$ /* 2-hop neighbors */
- 9: $N_3(v) \leftarrow \{w \mid \exists u \in N_2(v), w \in N(u), w \notin N_1(v) \cup N_2(v) \cup \{v\}\}$ /* 3-hop neighbors */

```

10:    $\mathcal{N}(v) \leftarrow N_1(v) \cup N_2(v) \cup N_3(v)$  /* All vertices within distance 3 */

11:
12:  /* Enumerate all 4-vertex sets containing  $v$  within  $\mathcal{N}(v) \cup \{v\}$  */

13:  for each triple  $\{a, b, c\} \subseteq \mathcal{N}(v)$  with  $a < b < c$  (lexicographically) do
14:     $S \leftarrow \{v, a, b, c\}$ 
15:    if  $G[S]$  is connected then
16:      Determine the isomorphism type  $\tau$  of  $G[S]$  among the 6 graphlet types
17:       $\phi(G)[\tau] \leftarrow \phi(G)[\tau] + 1$ 
18:    end if
19:  end for
20: end for
21:
22: /* Correct for overcounting: each graphlet counted once per vertex
   in it */
23:  $\phi(G) \leftarrow \phi(G)/4$ 
24: return  $\phi(G)$ 

```

Why this works:

- **Completeness:** Every connected 4-graphlet has diameter at most 3 (e.g., path P_4 has diameter 3). Therefore, all 4 vertices of any connected graphlet lie within distance 3 of at least one of its vertices. By considering all vertices v and their 3-hop neighborhoods, we find all connected 4-graphlets.
- **Correctness of counting:** Each 4-graphlet is counted exactly 4 times—once for each vertex in it (when that vertex is chosen as the "center" v). Dividing by 4 gives the correct count.

Running time analysis:

For each vertex $v \in V(G)$:

- Computing $N_1(v)$: $O(d)$
- Computing $N_2(v)$: $O(d \cdot d) = O(d^2)$ (at most d neighbors, each with degree $\leq d$)
- Computing $N_3(v)$: $O(d^2 \cdot d) = O(d^3)$
- $|\mathcal{N}(v)| \leq d + d^2 + d^3 = O(d^3)$
- Enumerating triples from $\mathcal{N}(v)$: $\binom{|\mathcal{N}(v)|}{3} = O(d^9)$

- Checking connectivity and isomorphism type of $G[S]$: $O(1)$ (4 vertices, constant time)

Cost per vertex: $O(d^9)$ (dominated by triple enumeration)

Total cost: $O(n \cdot d^9)$

Important note: While this is $O(d^9)$ per vertex, for **constant d **, this is $O(n)$ overall, which is polynomial in n .

Optimization to reduce the exponent:

We can improve the algorithm by being smarter about enumeration. Instead of considering all triples in $\mathcal{N}(v)$, we can:

1. **For each pair** $\{a, b\} \subseteq N_1(v)$ (the 1-hop neighbors):
 - Enumerate $c \in N_1(v) \cup N(a) \cup N(b)$
 - Check if $\{v, a, b, c\}$ is connected
 - Cost: $O(d^2 \cdot d) = O(d^3)$ per vertex
2. **For each** $a \in N_1(v)$ **and each pair** $\{b, c\} \subseteq N(a) \setminus \{v\}$:
 - Check if $\{v, a, b, c\}$ is connected
 - Cost: $O(d \cdot d^2) = O(d^3)$ per vertex

This improved algorithm achieves ** $O(n \cdot d^3)$ total time**.

Detailed optimized algorithm:

```

1: Initialize  $\phi(G) = [0, 0, 0, 0, 0, 0]$ 
2: Initialize hash table  $\mathcal{H}$  to track processed 4-sets
3:
4: for each vertex  $v \in V(G)$  do
5:   /* Case 1:  $v$  and two of its neighbors + one more */
6:   for each pair  $\{a, b\} \subseteq N(v)$  with  $a < b$  do
7:      $\mathcal{C} \leftarrow N(v) \cup N(a) \cup N(b)$ 
8:     for each  $c \in \mathcal{C}$  with  $c > \max(v, a, b)$  do
9:        $S \leftarrow \{v, a, b, c\}$  (in sorted order)
10:      if  $S \notin \mathcal{H}$  and  $G[S]$  is connected then
11:        Add  $S$  to  $\mathcal{H}$ 
12:        Identify isomorphism type  $\tau$  of  $G[S]$ 
13:         $\phi(G)[\tau] \leftarrow \phi(G)[\tau] + 1$ 
14:      end if
15:    end for

```

```

16: end for
17:
18: /* Case 2:  $v$ , one neighbor  $a$ , and two neighbors of  $a$  */
19: for each  $a \in N(v)$  do
20:   for each pair  $\{b, c\} \subseteq N(a) \setminus \{v\}$  with  $b < c$  do
21:      $S \leftarrow \{v, a, b, c\}$  (in sorted order)
22:     if  $S \notin \mathcal{H}$  and  $G[S]$  is connected then
23:       Add  $S$  to  $\mathcal{H}$ 
24:       Identify isomorphism type  $\tau$  of  $G[S]$ 
25:        $\phi(G)[\tau] \leftarrow \phi(G)[\tau] + 1$ 
26:     end if
27:   end for
28: end for
29: end for
30: return  $\phi(G)$ 

```

Analysis of optimized algorithm:

- Case 1: $O(d^2)$ pairs $\{a, b\}$, each considers $O(d)$ candidates for c : $O(d^3)$ per vertex
- Case 2: $O(d)$ neighbors a , each has $O(d^2)$ pairs: $O(d^3)$ per vertex
- Hash table operations: $O(1)$ expected time
- Isomorphism type identification: $O(1)$ for 4 vertices (check edge pattern)

Total time: $O(n \cdot d^3)$, which is **polynomial in n for constant d **.

Identifying isomorphism types efficiently: For a 4-vertex induced subgraph $G[S]$ with m edges:

- $m = 3$: Check structure (path vs star vs triangle+tail)
- $m = 4$: Must be C_4 (4-cycle)
- $m = 5$: Must be $K_4 - e$
- $m = 6$: Must be K_4

This can be done in $O(1)$ time by checking the degree sequence and specific edges.

7. **Problem 7:** Let G be a graph. Consider the following variant of the 1-WL, computing a coloring $D_t^1 : V(G) \rightarrow \mathbb{N}$, for $t > 0$, where

$$D_t^1(v) = \text{Recolor} \left(D_{(t-1)}^1(v), \{\{D_{(t-1)}^1(w) \mid w \in N(v)\}\}, \{\{D_{(t-1)}^1(w) \mid w \notin N(v)\}\} \right).$$

Show that the above variant of the 1-WL has the same expressivity as the ordinary 1-WL in distinguishing non-isomorphic graphs.

Ans: We need to show that this variant (which we call D-1-WL) and the ordinary 1-WL have the same expressivity, i.e., they distinguish exactly the same pairs of non-isomorphic graphs.

- Let $C_t^1(v)$ denote the color of vertex v after t iterations of the ordinary 1-WL.
- Let $D_t^1(v)$ denote the color of vertex v after t iterations of the D-1-WL variant.

Recall the ordinary 1-WL refinement:

$$C_t^1(v) = \text{Recolor}\left(C_{t-1}^1(v), \{\{C_{t-1}^1(w) \mid w \in N(v)\}\}\right).$$

The D-1-WL additionally considers the multiset of colors of **non-neighbors**:

$$D_t^1(v) = \text{Recolor}\left(D_{t-1}^1(v), \{\{D_{t-1}^1(w) \mid w \in N(v)\}\}, \{\{D_{t-1}^1(w) \mid w \notin N(v)\}\}\right).$$

Direction 1: Ordinary 1-WL \preceq D-1-WL

This direction is immediate since D-1-WL uses more information than ordinary 1-WL. Specifically:

- D-1-WL considers both the neighbor multiset and the non-neighbor multiset.
- Ordinary 1-WL only considers the neighbor multiset.

Therefore, if ordinary 1-WL can distinguish two graphs G and H , then D-1-WL can also distinguish them, since it has access to all the information that ordinary 1-WL uses (and more).

Formally: If $C_\infty^1(v) \neq C_\infty^1(w)$ for some vertices $v \in V(G)$ and $w \in V(H)$ that would otherwise correspond under an isomorphism, then the same reasoning applies to D_∞^1 , implying $D_\infty^1(v) \neq D_\infty^1(w)$.

Direction 2: D-1-WL \preceq Ordinary 1-WL

We need to show that the additional information (non-neighbor multiset) does not increase the distinguishing power.

Key observation: The non-neighbor multiset is redundant given the neighbor multiset and the total color distribution in the graph.

Why the non-neighbor information is redundant:

For a vertex v in a graph G with n vertices, let:

- $M_N(v) = \{\{D_{t-1}^1(w) \mid w \in N(v)\}\}$ be the neighbor multiset.
- $M_{\bar{N}}(v) = \{\{D_{t-1}^1(w) \mid w \notin N(v)\}\}$ be the non-neighbor multiset.
- $M_{\text{total}}(G) = \{\{D_{t-1}^1(w) \mid w \in V(G)\}\}$ be the total multiset of colors in G .

Then we have the fundamental relation:

$$M_{\text{total}}(G) = \{\{D_{t-1}^1(v)\}\} \uplus M_N(v) \uplus M_{\bar{N}}(v),$$

where \uplus denotes multiset union.

Therefore:

$$M_{\bar{N}}(v) = M_{\text{total}}(G) \setminus (\{\{D_{t-1}^1(v)\}\} \uplus M_N(v)),$$

where \setminus denotes multiset difference.

This shows that $M_{\bar{N}}(v)$ is completely determined by:

1. The total color distribution $M_{\text{total}}(G)$ (global information)
2. The vertex's own color $D_{t-1}^1(v)$ (local information)
3. The neighbor multiset $M_N(v)$ (local information)

Simulating D-1-WL with ordinary 1-WL:

We can simulate the D-1-WL using a modified version of ordinary 1-WL that includes global color histograms.

Define an augmented 1-WL variant where the color refinement is:

$$\tilde{C}_t^1(v) = \text{Recolor} \left(\tilde{C}_{t-1}^1(v), \{\{\tilde{C}_{t-1}^1(w) \mid w \in N(v)\}\}, M_{\text{total}}(G, t-1) \right),$$

where $M_{\text{total}}(G, t-1)$ is the color histogram of G at iteration $t-1$.

Claim: This augmented variant has the same distinguishing power as ordinary 1-WL.

[Proof] The ordinary 1-WL on the disjoint union $G \sqcup H$ implicitly captures the global color distributions. When we run ordinary 1-WL on $G \sqcup H$:

- Two vertices $v \in V(G)$ and $w \in V(H)$ can only get the same stable color if they have the same neighborhood structure.

- The stable partition of $V(G) \sqcup V(H)$ ensures that if G and H have different color distributions at any level, this will eventually be reflected in different stable colors.

More precisely, if ordinary 1-WL distinguishes G and H , it means either:

1. The color histograms differ: $|\{v \in V(G) \mid C_\infty^1(v) = c\}| \neq |\{w \in V(H) \mid C_\infty^1(w) = c\}|$ for some color c , or
2. There exist $v \in V(G)$ and $w \in V(H)$ with the same initial label but different stable colors due to neighborhood differences.

In both cases, the D-1-WL will also distinguish them, because:

- If the color histograms differ, then $M_{\text{total}}(G) \neq M_{\text{total}}(H)$.
- If neighborhood structures differ for some vertices, the neighbor multisets already capture this (the non-neighbor information is redundant as shown above).

Formal argument using the redundancy:

Suppose D-1-WL distinguishes two graphs G and H , but ordinary 1-WL does not. This means:

- Ordinary 1-WL: The stable color histograms of G and H are identical, i.e., $\{\{C_\infty^1(v) \mid v \in V(G)\}\} = \{\{C_\infty^1(w) \mid w \in V(H)\}\}$.
- D-1-WL: The stable color histograms differ, or vertices with matching ordinary 1-WL colors get different D-1-WL colors.

But if ordinary 1-WL assigns the same color histograms, then:

$$M_{\text{total}}(G, t) = M_{\text{total}}(H, t) \text{ for all } t \geq 0.$$

Now, for any pair of vertices $v \in V(G)$ and $w \in V(H)$ with $C_{t-1}^1(v) = C_{t-1}^1(w)$, we also have:

- $\{\{C_{t-1}^1(u) \mid u \in N(v)\}\} = \{\{C_{t-1}^1(u') \mid u' \in N(w)\}\}$ (by definition of color refinement).
- Since $M_{\text{total}}(G, t-1) = M_{\text{total}}(H, t-1)$ and the neighbor multisets are equal, we must also have:

$$\{\{C_{t-1}^1(u) \mid u \notin N(v)\}\} = \{\{C_{t-1}^1(u') \mid u' \notin N(w)\}\}.$$

Therefore, the non-neighbor multisets must also be equal, which means $D_t^1(v) = D_t^1(w)$.

This contradicts the assumption that D-1-WL distinguishes G and H while ordinary 1-WL does not.

Conclusion. We have shown both directions:

- Ordinary 1-WL \preceq D-1-WL (trivial, since D-1-WL uses more information)
- D-1-WL \preceq Ordinary 1-WL (non-trivial, using the redundancy argument)

Hence,

$$\text{Ordinary 1-WL} \equiv \text{D-1-WL.}$$



8. **Problem 8:** Let X be a set, and let $k_1 : X \times X \rightarrow \mathbb{R}$ and $k_2 : X \times X \rightarrow \mathbb{R}$ be two kernels with feature maps $\phi_1 : X \rightarrow \mathbb{R}^d$ and $\phi_2 : X \rightarrow \mathbb{R}^e$, respectively. Show that $k_+(x, y) = k_1(x, y)^2 + k_2(x, y)^3$ for $x, y \in X$ is a valid positive semidefinite kernel.

Ans: We have to prove that k_+ is a valid kernel by proving it is positive semidefinite, i.e., for any $n \geq 1$ and any $x_1, \dots, x_n \in X$, the Gram matrix K_+ with entries $(K_+)_ij = k_+(x_i, x_j)$ is positive semidefinite.

We use the fact that kernels are closed under certain operations:

- Addition: If k and k' are kernels, then $k + k'$ is a kernel.
- Multiplication: If k and k' are kernels, then $k \cdot k'$ is a kernel (pointwise product).
- Powers: If k is a kernel, then k^n (product of k with itself n times) is a kernel.

Step 1: Show k_1^2 is a kernel.

Since k_1 is a kernel, we can write $k_1(x, y) = \langle \phi_1(x), \phi_1(y) \rangle$.

The pointwise product of two kernels is a kernel. Therefore:

$$k_1^2(x, y) = k_1(x, y) \cdot k_1(x, y)$$

is a kernel.

Explicit feature map: We can construct a feature map for k_1^2 using the tensor product:

$$\psi_1(x) = \phi_1(x) \otimes \phi_1(x) \in \mathbb{R}^{d \times d}.$$

Then:

$$k_1^2(x, y) = \langle \phi_1(x), \phi_1(y) \rangle^2 = \langle \phi_1(x) \otimes \phi_1(x), \phi_1(y) \otimes \phi_1(y) \rangle = \langle \psi_1(x), \psi_1(y) \rangle.$$

Step 2: Show k_2^3 is a kernel.

Similarly, since k_2 is a kernel with $k_2(x, y) = \langle \phi_2(x), \phi_2(y) \rangle$:

$$k_2^3(x, y) = k_2(x, y) \cdot k_2(x, y) \cdot k_2(x, y)$$

is a kernel (product of three kernels).

Explicit feature map: Using the tensor product:

$$\psi_2(x) = \phi_2(x) \otimes \phi_2(x) \otimes \phi_2(x) \in \mathbb{R}^{e \times e \times e}.$$

Then:

$$k_2^3(x, y) = \langle \phi_2(x), \phi_2(y) \rangle^3 = \langle \psi_2(x), \psi_2(y) \rangle.$$

Step 3: Show $k_+ = k_1^2 + k_2^3$ is a kernel.

Since kernels are closed under addition, and we have shown that k_1^2 and k_2^3 are both kernels:

$$k_+(x, y) = k_1^2(x, y) + k_2^3(x, y)$$

is a kernel.

Explicit feature map: We can construct a feature map for k_+ by concatenating:

$$\phi_+(x) = [\psi_1(x), \psi_2(x)] \in \mathbb{R}^{d^2 + e^3}.$$

Then:

$$k_+(x, y) = \langle \psi_1(x), \psi_1(y) \rangle + \langle \psi_2(x), \psi_2(y) \rangle = \langle \phi_+(x), \phi_+(y) \rangle.$$

Verification of positive semidefiniteness:

For any n points $x_1, \dots, x_n \in X$ and any coefficients $c_1, \dots, c_n \in \mathbb{R}$:

$$\begin{aligned}
\sum_{i,j=1}^n c_i c_j k_+(x_i, x_j) &= \sum_{i,j=1}^n c_i c_j [k_1(x_i, x_j)^2 + k_2(x_i, x_j)^3] \\
&= \sum_{i,j=1}^n c_i c_j k_1(x_i, x_j)^2 + \sum_{i,j=1}^n c_i c_j k_2(x_i, x_j)^3 \\
&= \sum_{i,j=1}^n c_i c_j \langle \psi_1(x_i), \psi_1(x_j) \rangle + \sum_{i,j=1}^n c_i c_j \langle \psi_2(x_i), \psi_2(x_j) \rangle \\
&= \left\langle \sum_{i=1}^n c_i \psi_1(x_i), \sum_{j=1}^n c_j \psi_1(x_j) \right\rangle + \left\langle \sum_{i=1}^n c_i \psi_2(x_i), \sum_{j=1}^n c_j \psi_2(x_j) \right\rangle \\
&= \left\| \sum_{i=1}^n c_i \psi_1(x_i) \right\|^2 + \left\| \sum_{i=1}^n c_i \psi_2(x_i) \right\|^2 \\
&\geq 0.
\end{aligned}$$

Therefore, the Gram matrix is positive semidefinite, confirming that k_+ is a valid kernel.

Conclusion. The function $k_+(x, y) = k_1(x, y)^2 + k_2(x, y)^3$ is a valid positive semidefinite kernel, with feature map $\phi_+(x) = [\phi_1(x) \otimes \phi_1(x), \phi_2(x) \otimes \phi_2(x)]$.

